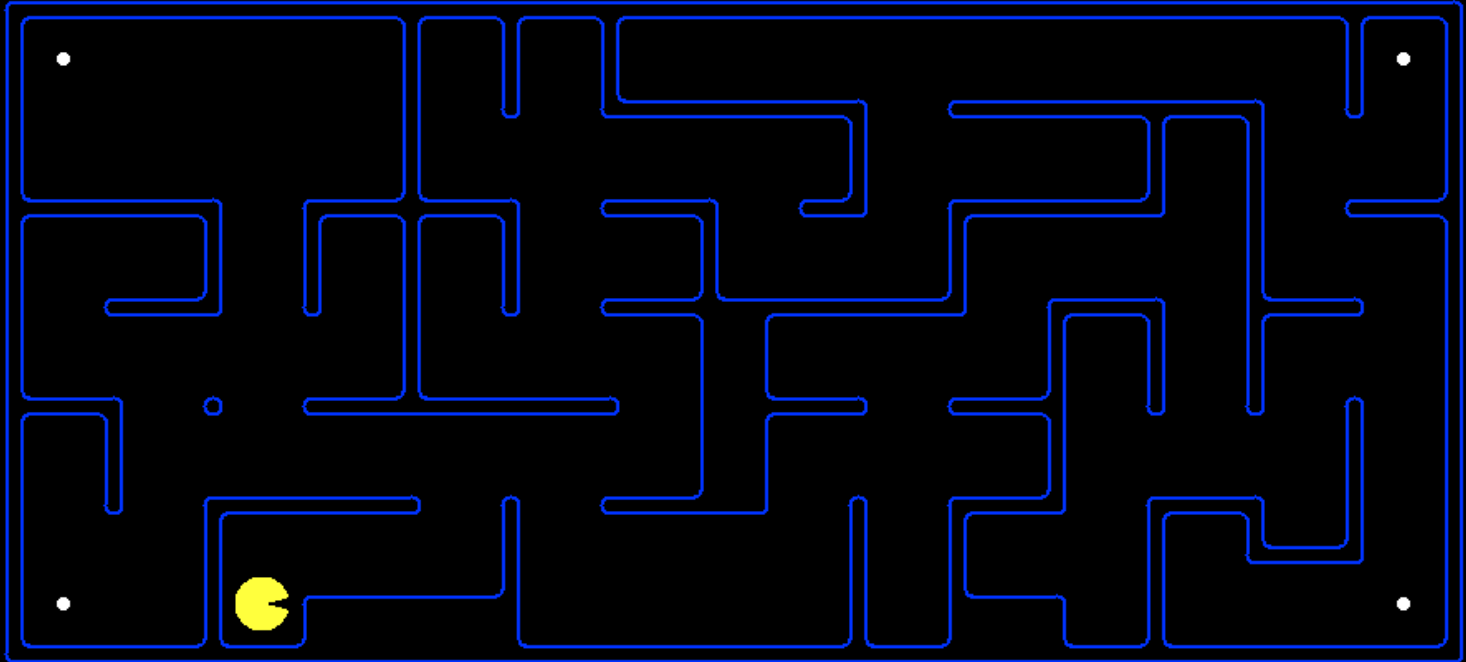


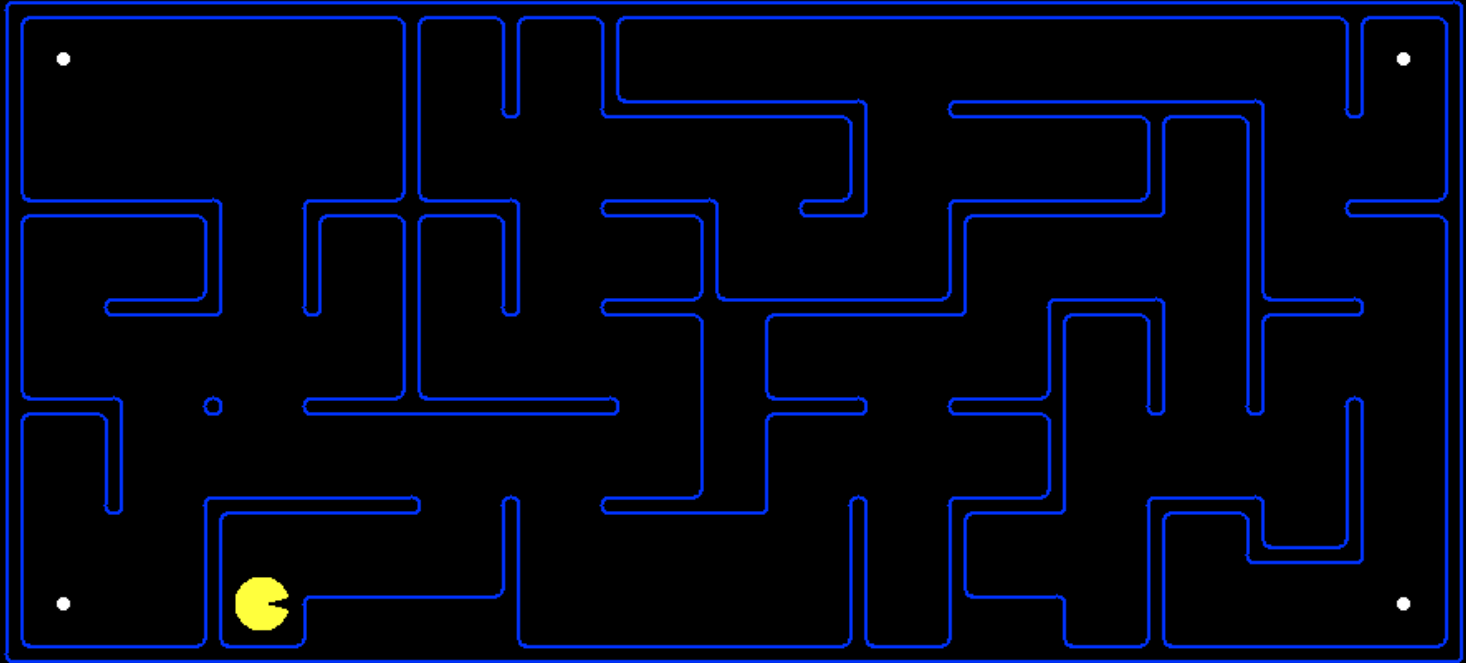
# Introduction to Artificial Intelligence

Lecture 3: Constraint satisfaction problems


Prof. Gilles Louppe  
[g.louppe@uliege.be](mailto:g.louppe@uliege.be)



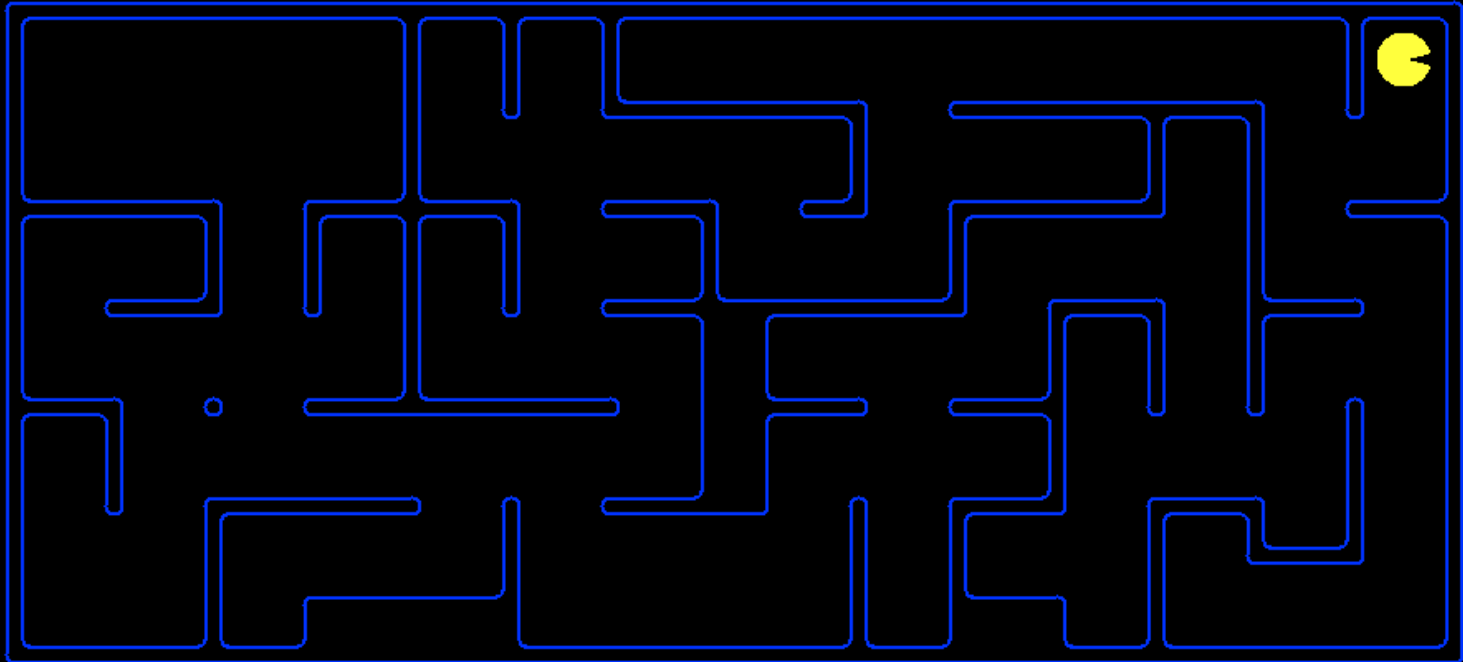
Hmmm, let me think...



(...)



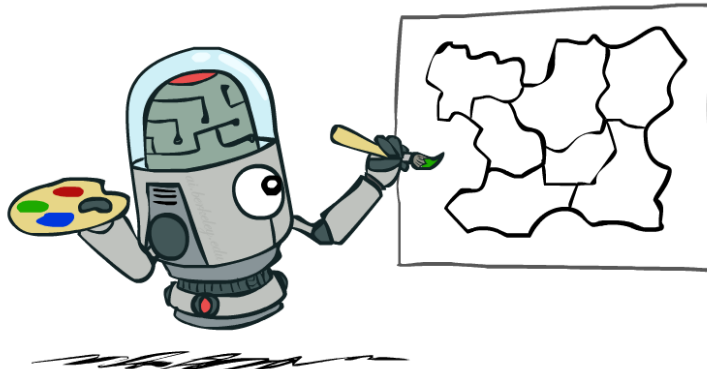
(some time later)



Solution found! [Can we do better?]

# Today

- **Constraint satisfaction problems:**
  - Exploiting the representation of a state to accelerate search.
  - Backtracking.
  - Generic heuristics.
- **Logical agents**
  - Propositional logic for reasoning about the world.
  - ... and its connection with CSPs.



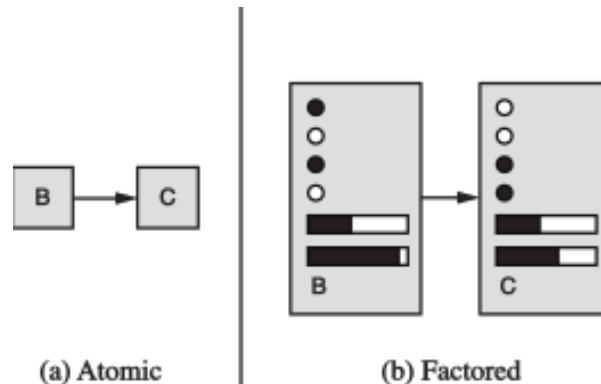
# Constraint satisfaction problems

# Motivation

In standard search problems:

- States are evaluated by domain-specific heuristics.
- States are tested by a domain-specific function to determine if the goal is achieved.
- From the point of view of the search algorithms however, **states are atomic**.

Instead, if states have **a factored representation**, then the structure of states can be exploited to improve the **efficiency of the search**.





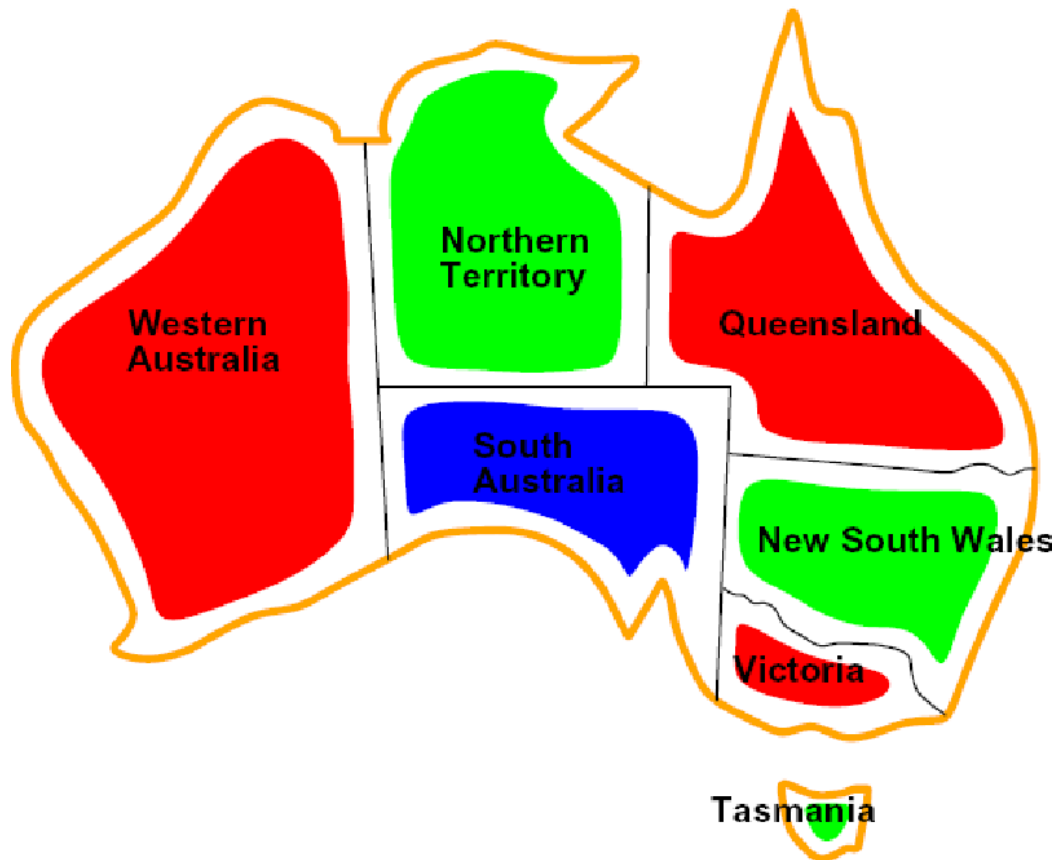
# Constraint satisfaction problems

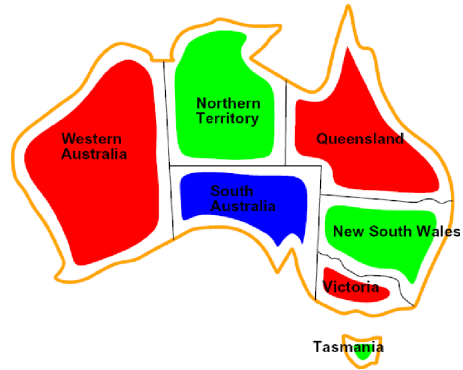
- **Constraint satisfaction problem** solvers take advantage of factored state representations and use **general-purpose** heuristics to solve complex problems.
- CSPs are specialized to a family of search sub-problems.
- Main idea: eliminate large portions of the search space all at once, by identifying combinations of variable/value that violate constraints.

Formally, a **constraint satisfaction problem** (CSP) consists of three components  $X$ ,  $D$  and  $C$ :

- $X$  is a set of **variables**,  $\{X_1, \dots, X_n\}$ ,
- $D$  is a set of **domains**,  $\{D_1, \dots, D_n\}$ , one for each variable,
- $C$  is a set of **constraints** that specify allowable combinations of values.

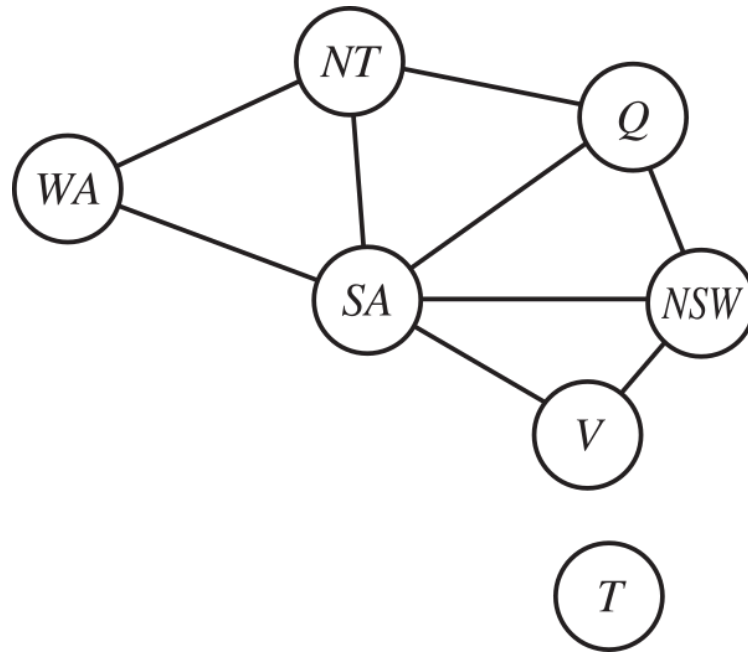
## Example: Map coloring





- Variables:  $X = \{WA, NT, Q, NSW, V, SA, T\}$
- Domains:  $D_i = \{\text{red, green, blue}\}$  for each variable.
- Constraints:  $C = \{SA \neq WA, SA \neq NT, SA \neq Q, \dots\}$ 
  - Implicit:  $WA \neq NT$
  - Explicit:  $(WA, NT) \in \{\{\text{red, green}\}, \{\text{red, blue}\}, \dots\}$
- Solutions are **assignments** of values to the variables such that constraints are all satisfied.
  - e.g.,  $\{WA = \text{red}, NT = \text{green}, Q = \text{red}, SA = \text{blue}, NSW = \text{green}, V = \text{red}, T = \text{green}\}$

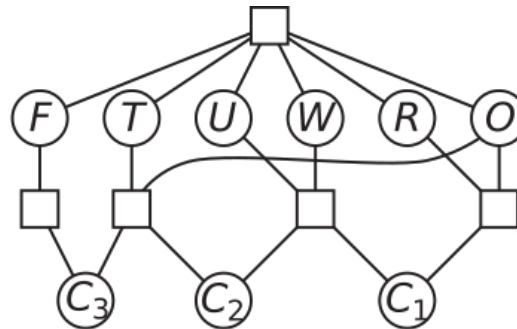
# Constraint (hyper)graph



- **Nodes** = variables of the problems
- **Edges** = constraints in the problem involving the variables associated to the end nodes.
- General purpose CSP algorithms **use the graph structure** to speedup search.
  - e.g., Tasmania is an independent subproblem.

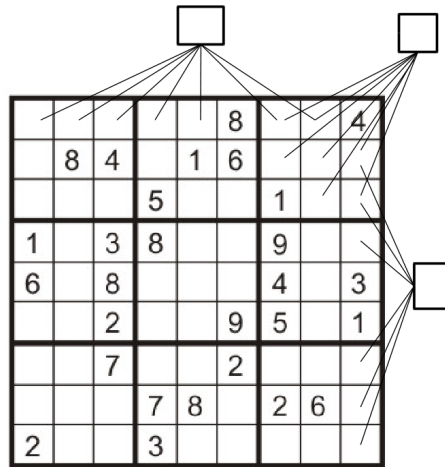
## Example: Cryptarithmic

$$\begin{array}{r} T W O \\ + T W O \\ \hline F O U R \end{array}$$




- Variables:  $\{T, W, O, F, U, R, C_1, C_2, C_3\}$
- Domains:  $D_i = \{0, 1, 2, 3, 4, 5, 6, 7, 8, 9\}$
- Constraints:
  - $\text{alldiff}(T, W, O, F, U, R)$
  - $O + O = R + 10 \times C_1$
  - $C_1 + W + W = U + 10 \times C_2$
  - ...

## Example: Sudoku



- Variables: each (open) square
- Domains:  $D_i = \{1, 2, 3, 4, 5, 6, 7, 8, 9\}$
- Constraints:
  - 9-way **alldiff** for each column
  - 9-way **alldiff** for each row
  - 9-way **alldiff** for each region




## Example: The Waltz algorithm

Procedure for interpreting 2D line drawings of solid polyhedra as 3D objects.  
Early example of an AI computation posed as a CSP.

CSP formulation:

- Each intersection is a variable.
- Adjacent intersections impose constraints on each other.
- Solutions are physically realizable 3D objects.





# Variations on the CSP formalism

- Discrete variables

- Finite domains

- Size  $d$  means  $O(d^n)$  complete assignments.
    - e.g., boolean CSPs, including the SAT boolean satisfiability problem (NP-complete).

- Infinite domains

- e.g., job scheduling, variables are start/end days for for each job.
    - need a constraint language, e.g.  $start_1 + 5 \leq start_2$ .
    - Solvable for linear constraints, undecidable otherwise.

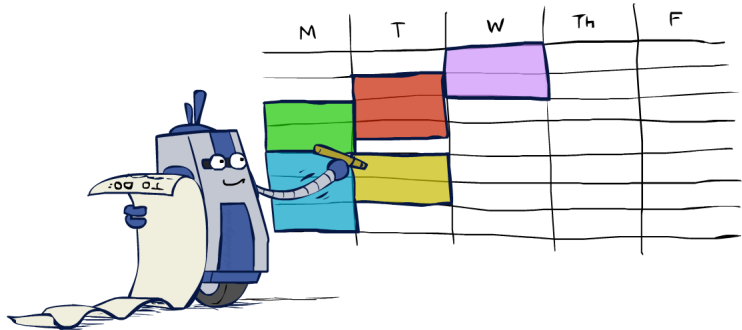
- Continuous variables

- e.g., precise start/end times of experiments.
  - Linear constraints solvable in polynomial time by LP methods.

- Varieties of constraints
  - Unary constraint involve a single variable.
    - Equivalent to reducing the domain, e.g.  $SA \neq \text{green}$ .
  - Binary constraints involve pairs of variables, e.g.  $SA \neq WA$ .
  - Higher-order constraints involve 3 or more variables.
- Preferences (soft constraints)
  - e.g., red is better than green.
  - Often representable by a cost for each variable assignment.
  - Results in constraint optimization problems.
  - (We will ignore those in this course.)

# Real-world examples

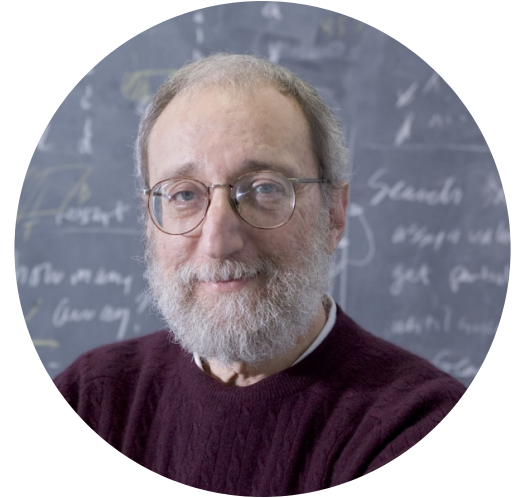
- Assignment problems
  - e.g., who teaches what class
- Timetabling problems
  - e.g., which class is offered when and where?
- Hardware configuration
- Spreadsheets
- Transportation scheduling
- Factory scheduling
- Circuit layout
- ... and many more



Notice that many real-world problems involve real-valued variables.

# Constraint programming

*Constraint Programming represents one of the closest approaches computer science has yet made to the Holy Grail of programming: the user states the problem, the computer solves it.*  
(Eugene Freuder)



Constraint programming is a programming paradigm in which the user specifies the program as a CSP. The resolution of the problem is left to the computer.


Examples:

- Prolog
- ECLiPSe

# Solving CSPs

# Standard search formulation

- CSPs can be cast as standard search problems.
  - For which we have solvers, including DFS, BFS or A\*.
- States are **partial assignments**:
  - The initial state is the empty assignment  $\{\}$ .
  - Actions: assign a value to an unassigned variable.
  - Goal test: the current assignment is complete and satisfies all constraints.
- This algorithm is **the same** for all CSPs!



What would BFS or DFS do? What problems does naive search have?

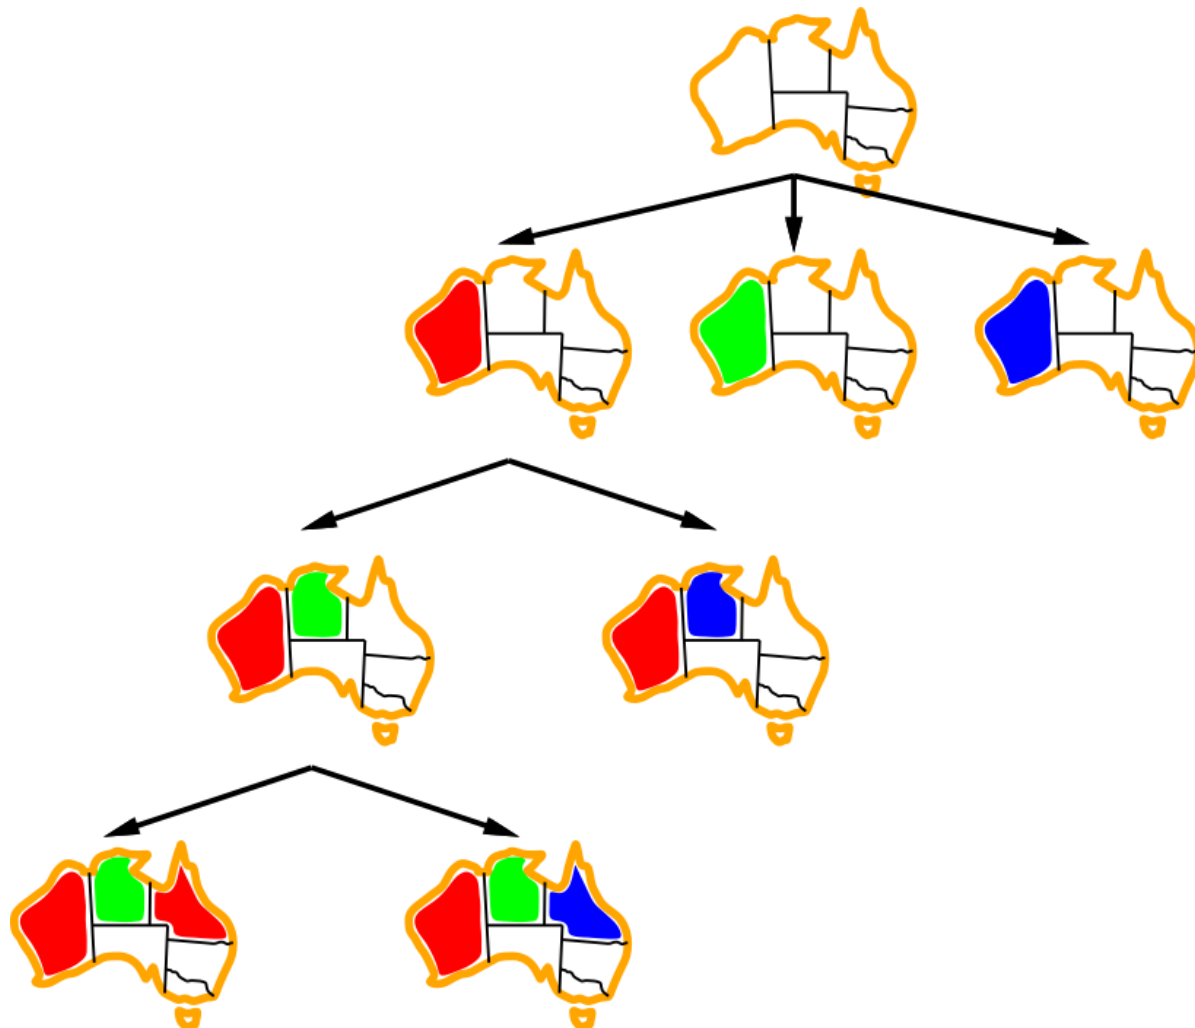
For  $n$  variables of domain size  $d$ :

- $b = (n - l)d$  at depth  $l$ ;
- we generate a tree with  $n!d^n$  leaves even if there are only  $d^n$  possible assignments!

# Backtracking search

- Backtracking search is a canonical uninformed algorithm for solving CSPs.
- Idea 1: **One variable at a time:**
  - The naive application of search algorithms ignores a crucial property: variable assignments are **commutative**. Therefore, fix the ordering.
    - $WA = \text{red}$  then  $NT = \text{green}$  is the same as  $NT = \text{green}$  then  $WA = \text{red}$ .
  - One only needs to consider assignments to a single variable at each step.
    - $b = d$  and there are  $d^n$  leaves.
- Idea 2: **Check constraints as you go:**
  - Consider only values which do not conflict with current partial assignment.
  - Incremental goal test.





```

function BACKTRACKING-SEARCH(csp) returns a solution, or failure
    return BACKTRACK({ }, csp)

function BACKTRACK(assignment, csp) returns a solution, or failure
    if assignment is complete then return assignment
    var  $\leftarrow$  SELECT-UNASSIGNED-VARIABLE(csp)
    for each value in ORDER-DOMAIN-VALUES(var, assignment, csp) do
        if value is consistent with assignment then
            add {var = value} to assignment
            inferences  $\leftarrow$  INFERENCE(csp, var, value)
            if inferences  $\neq$  failure then
                add inferences to assignment
                result  $\leftarrow$  BACKTRACK(assignment, csp)
                if result  $\neq$  failure then
                    return result
            remove {var = value} and inferences from assignment
    return failure

```

---

**Figure 6.5** A simple backtracking algorithm for constraint satisfaction problems. The algorithm is modeled on the recursive depth-first search of Chapter 3. By varying the functions SELECT-UNASSIGNED-VARIABLE and ORDER-DOMAIN-VALUES, we can implement the general-purpose heuristics discussed in the text. The function INFERENCE can optionally be used to impose arc-, path-, or  $k$ -consistency, as desired. If a value choice leads to failure (noticed either by INFERENCE or by BACKTRACK), then value assignments (including those made by INFERENCE) are removed from the current assignment and a new value is tried.

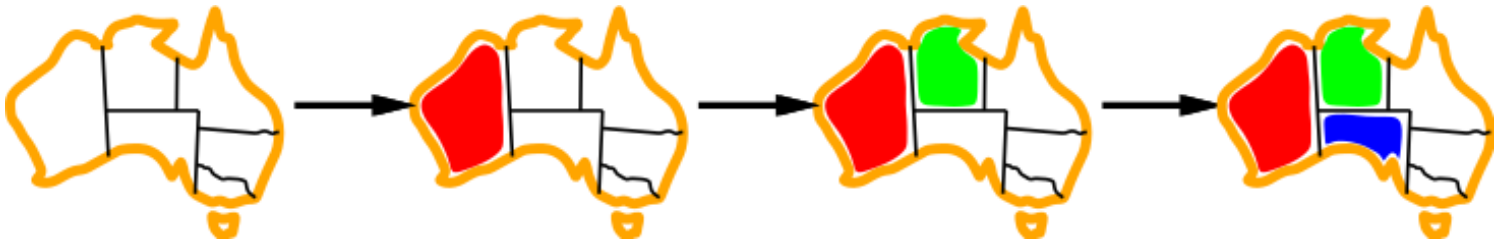
## Improving backtracking

Can we improve backtracking using **general-purpose** ideas, without domain-specific knowledge?

- **Ordering:**
  - Which variable should be assigned next?
  - In what order should its values be tried?
- **Filtering:** can we detect inevitable failure early?
- **Structure:** can we exploit the problem structure?

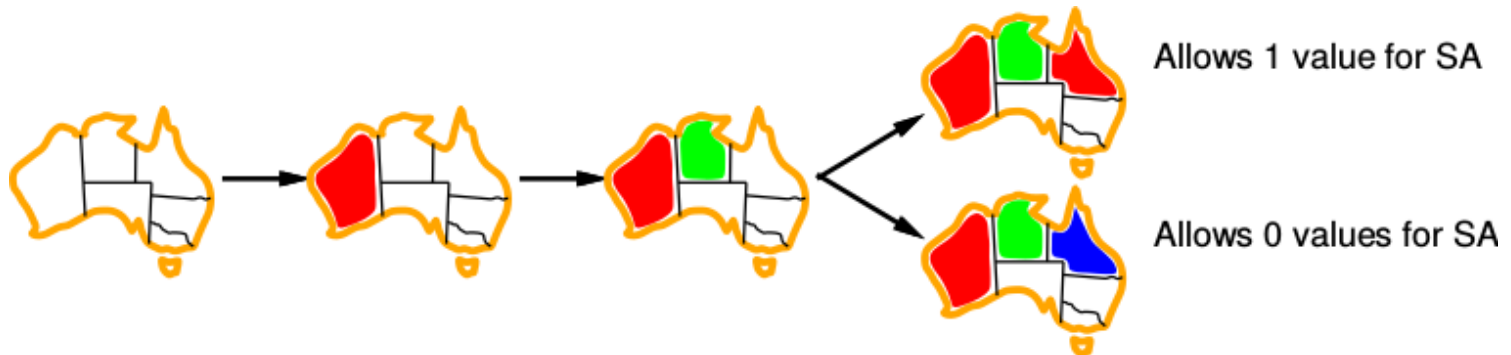
## Variable ordering

- **Minimum remaining values:** Choose the variable with the fewest legal values left in its domain.
- Also known as the **fail-first** heuristic.
  - Detecting failures quickly is equivalent to pruning large parts of the search tree.



## Value ordering

- **Least constraining value:** Given a choice of variable, choose the **least constraining value**.
- i.e., the value that rules out the fewest values in the remaining variables.



### Exercise

Why should variable selection be fail-first but value selection be fail-last?

## Filtering: Forward checking

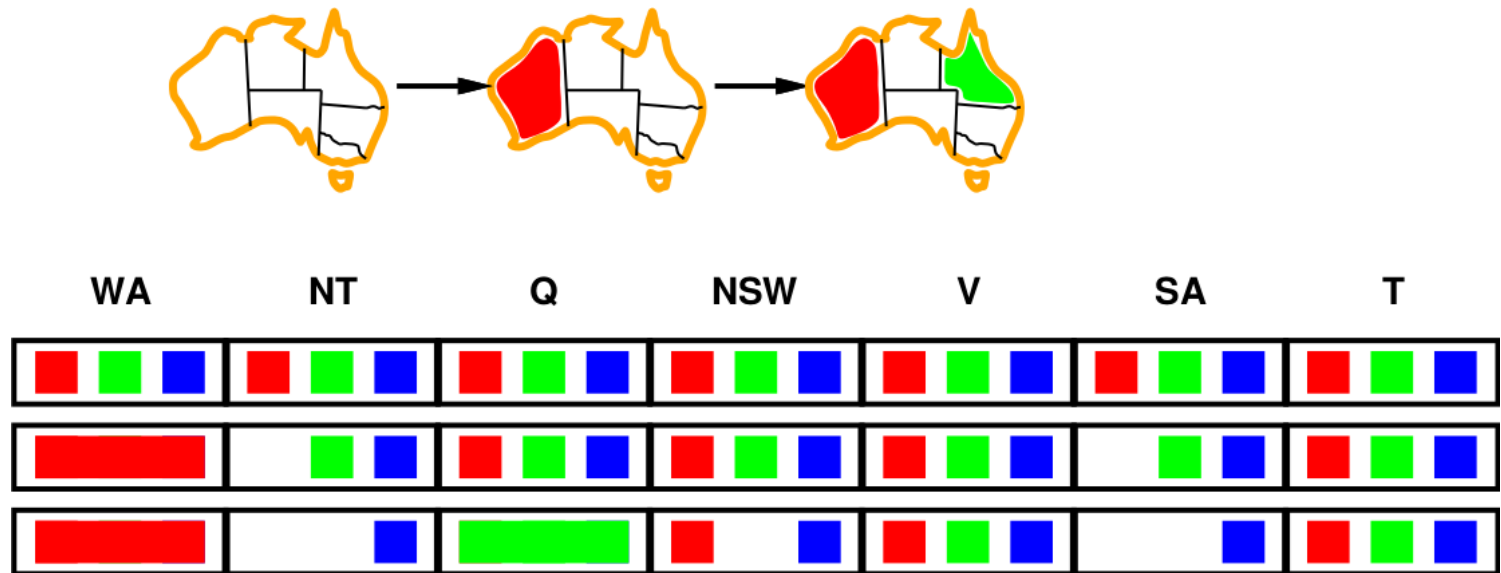
- Keep track of remaining legal values for unassigned variables.
  - Whenever a variable  $X$  is assigned, and for each unassigned variable  $Y$  that is connected to  $X$  by a constraint, delete from  $Y$ 's domain any value that is inconsistent.
- Terminate search when any variable has no legal value left.



WA	NT	Q	NSW	V	SA	T
<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>
<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>
<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>
<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>	<div><div></div><div></div><div></div></div>

## Filtering: Constraint propagation

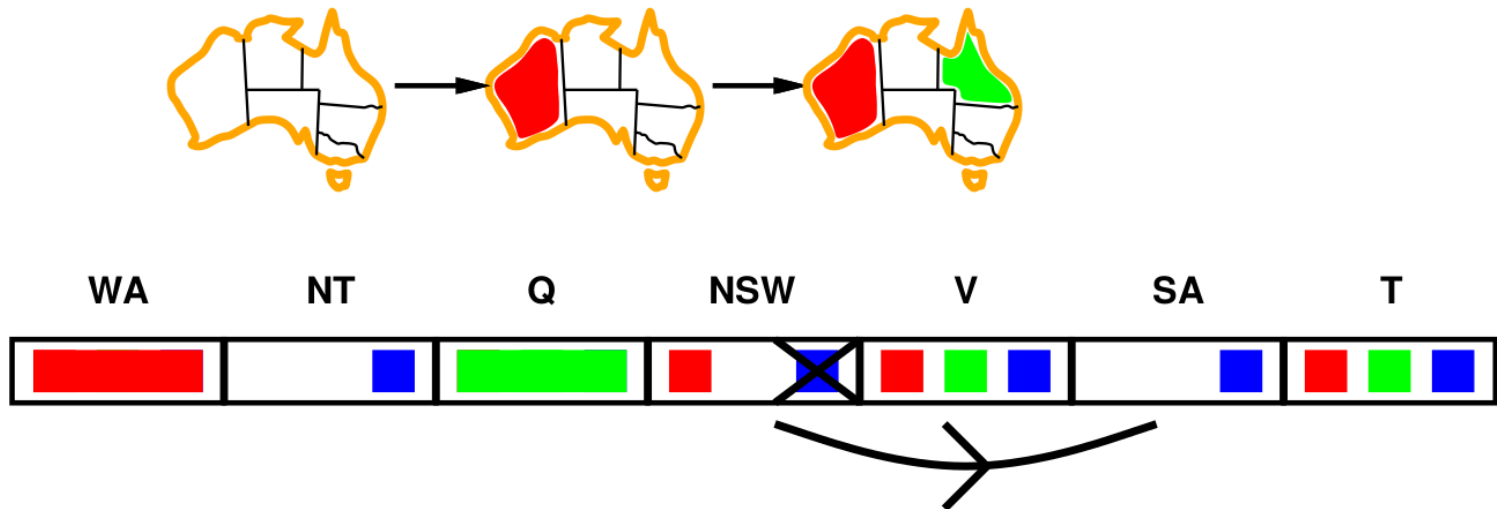
Forward checking propagates information assigned to unassigned variables, but does not provide early detection for all failures:



- *NT* and *SA* cannot both be blue!
- **Constraint propagation** repeatedly enforces constraints locally.

## Arc consistency

- An arc  $X \rightarrow Y$  is **consistent** if and only if for every value  $x$  in the domain of  $X$  there is some value  $y$  in the domain of  $Y$  that satisfies the associated binary constraint.
- Forward checking  $\Leftrightarrow$  enforcing consistency of arcs pointing to each new assignment.
- This principle can be generalized to enforce consistency for **all** arcs.





**function** AC-3(*csp*) **returns** false if an inconsistency is found and true otherwise

**inputs:** *csp*, a binary CSP with components ( $X$ ,  $D$ ,  $C$ )

**local variables:** *queue*, a queue of arcs, initially all the arcs in *csp*

**while** *queue* is not empty **do**

$(X_i, X_j) \leftarrow \text{REMOVE-FIRST}(\textit{queue})$

**if** REVISE(*csp*,  $X_i$ ,  $X_j$ ) **then**

**if** size of  $D_i = 0$  **then return** *false*

**for each**  $X_k$  **in**  $X_i.\text{NEIGHBORS} - \{X_j\}$  **do**

            add  $(X_k, X_i)$  to *queue*

**return** *true*

---

**function** REVISE(*csp*,  $X_i$ ,  $X_j$ ) **returns** true iff we revise the domain of  $X_i$

*revised*  $\leftarrow$  *false*

**for each**  $x$  **in**  $D_i$  **do**

**if** no value  $y$  in  $D_j$  allows  $(x,y)$  to satisfy the constraint between  $X_i$  and  $X_j$  **then**

            delete  $x$  from  $D_i$

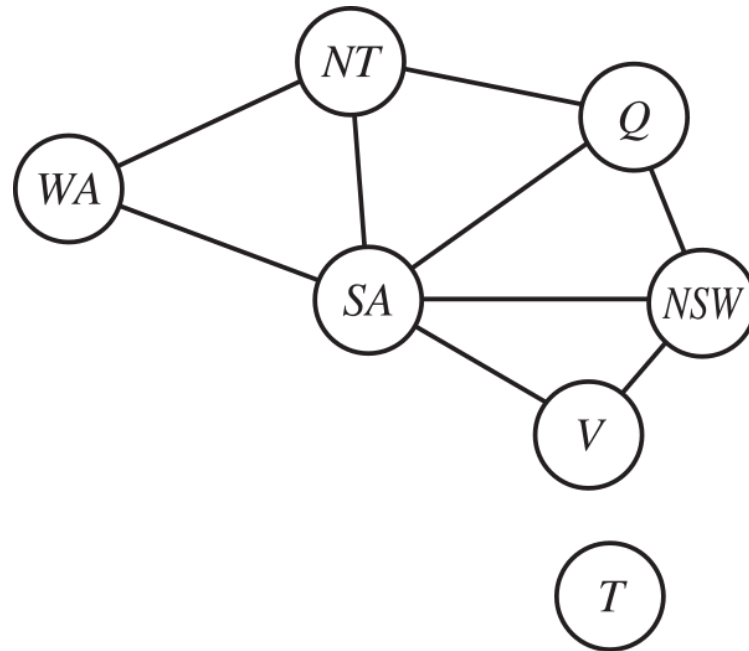
*revised*  $\leftarrow$  *true*

**return** *revised*

## Exercise

When in backtracking shall this procedure be called?

# Structure



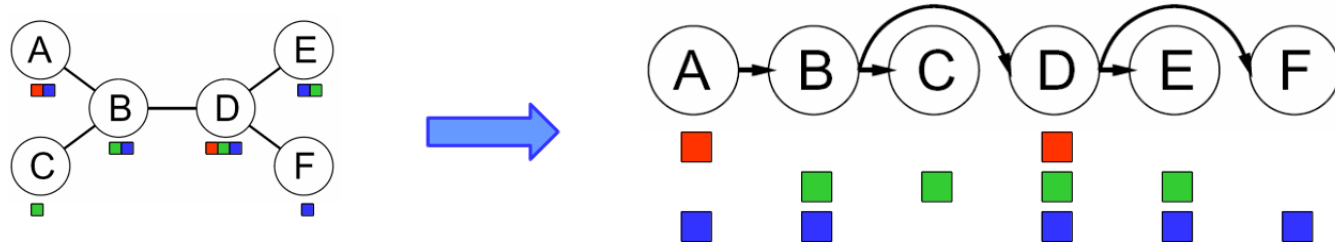
- Tasmania and mainland are **independent subproblems**.
  - Any solution for the mainland combined with any solution for Tasmania yields a solution for the whole map.
- Independence can be ascertained by finding **connected components** of the constraint graph.

## Time complexity

Assume each subproblem has  $c$  variables out of  $n$  in total. Then  $O\left(\frac{n}{c}d^c\right)$ .

- E.g.,  $n = 80, d = 2, c = 20$ .
- $2^{80} = 4$  billion years at 10 million nodes/sec.
- $4 \times 2^{20} = 0.4$  seconds at 10 million nodes/sec.

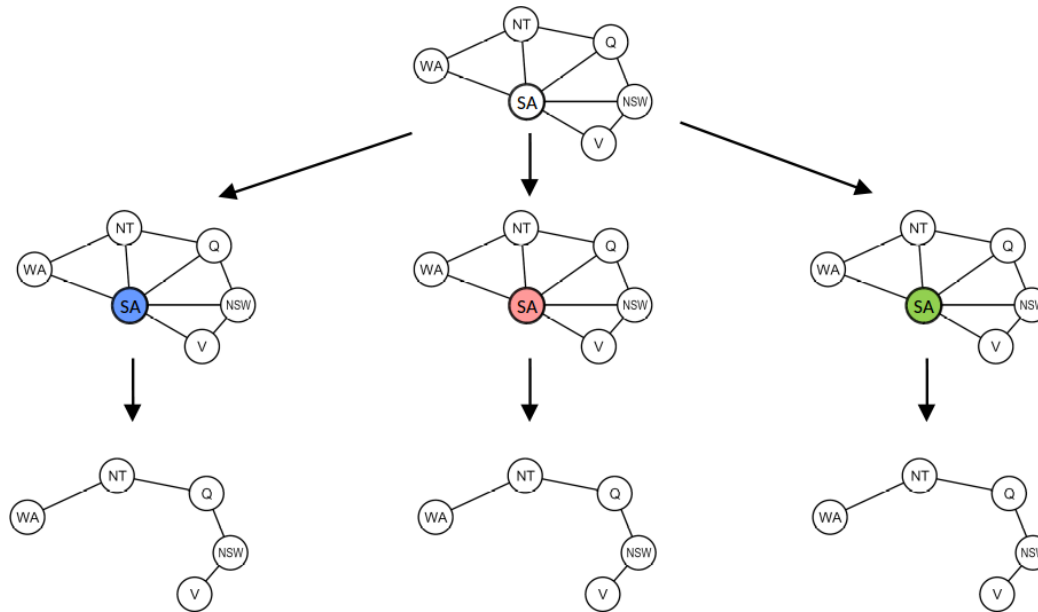
## Tree-structured CSPs



- Algorithm for tree-structured CSPs:
  - Order: choose a root variable, order variables so that parents precede children (topological sort).
  - Remove backward:
    - for  $i = n$  down to  $2$ , enforce arc consistency of  $parent(X_i) \rightarrow X_i$ .
  - Assign forward:
    - for  $i = 1$  to  $n$ , assign  $X_i$  consistently with its  $parent(X_i)$ .
- Time complexity:  $O(nd^2)$ 
  - Compare to general CSPs, where worst-case time is  $O(d^n)$ .

## Nearly tree-structured CSPs

- **Conditioning**: instantiate a variable, prune its neighbors' domains.
- **Cutset conditioning**:
  - Assign (in all ways) a set  $S$  of variables such that the remaining constraint graph is a tree.
  - Solve the residual CSPs (tree-structured).
  - If the residual CSP has a solution, return it together with the assignment for  $S$ .



# Logical agents

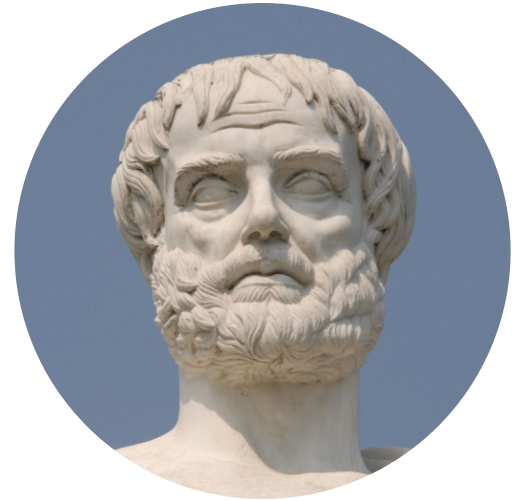
# The logicist tradition

- The rational thinking approach to artificial intelligence is concerned with the study of **irrefutable reasoning processes**. It ensures that all actions performed by an agent are formally **provable** from inputs and prior knowledge.
- The Greek philosopher Aristotle was one of the first to attempt to formalize rational thinking. His **syllogisms** provided a pattern for argument structures that always yield correct conclusion when given correct premises.

*All men are mortal.*

*Socrates is a man.*

*Therefore, Socrates is mortal.*

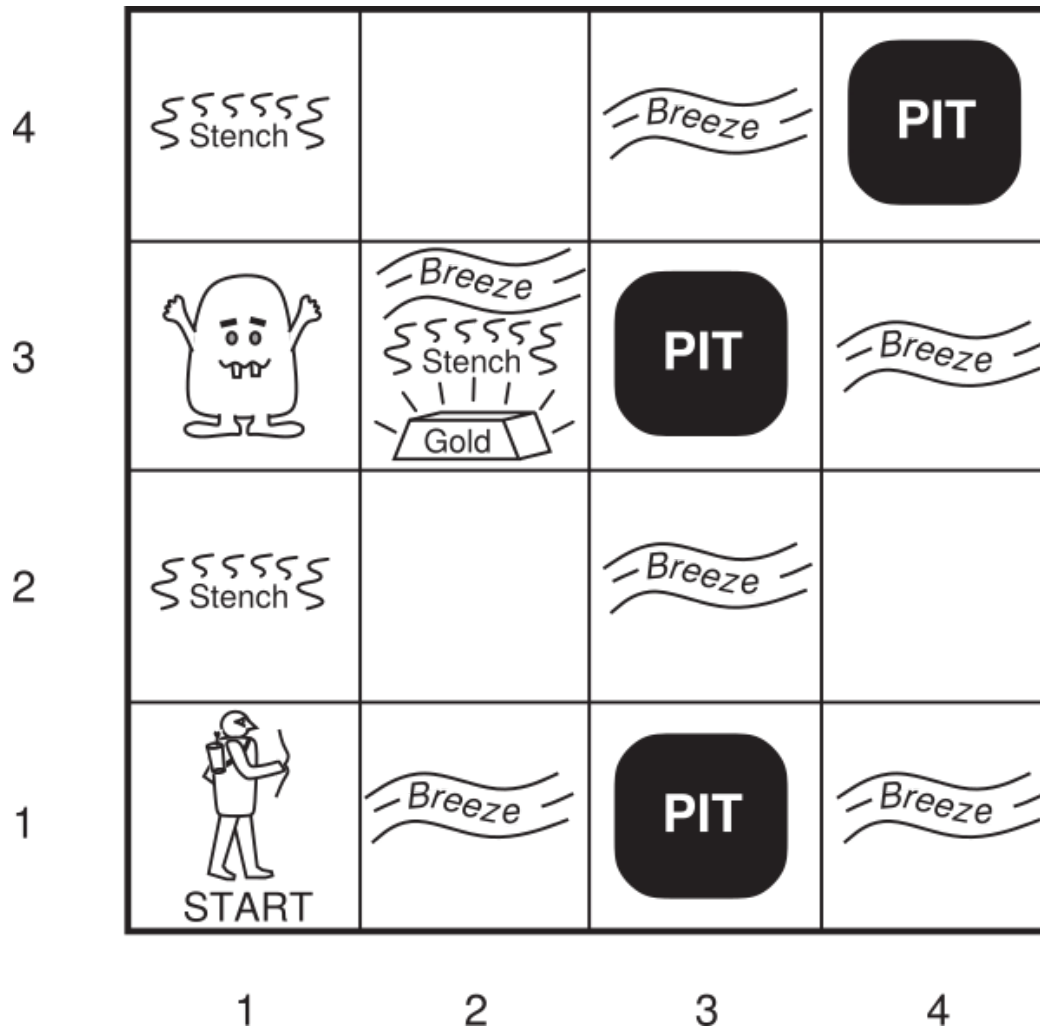


*(Aristotle, 384-322 BC)*

- Logicians of the 19th century developed a precise notation for statements about all kinds of objects in the world and relationships among them.
- By 1965, programs existed that could, in principle, solve any solvable problem described in logical notation.
- The logicist tradition within AI hopes to build on such programs to create intelligent systems.



# The Wumpus world



## PEAS description

- **Performance measure:**
  - +1000 for climbing out of the cave with gold;
  - -1000 for falling into a pit or being eaten by the wumpus;
  - -1 per step.
- **Environment:**
  - $4 \times 4$  grid of rooms;
  - The agent starts in the lower left square labeled  $[1, 1]$ , facing right;
  - Locations for gold, the wumpus and pits are chosen randomly from squares other than the start square.
- **Actuators:**
  - Forward, Turn left by  $90^\circ$  or Turn right by  $90^\circ$ .
- **Sensors:**
  - Squares adjacent to wumpus are **smelly**;
  - Squares adjacent to pit are **breezy**;
  - **Glitter** if gold is in the same square;
    - Gold is picked up by reflex, and cannot be dropped.
  - You **bump** if you walk into a wall.
  - The agent program receives the percept  $[\text{Stench}, \text{Breeze}, \text{Glitter}, \text{Bump}]$ .

## Wumpus world characterization

- **Deterministic**: Yes, outcomes are exactly specified.
- **Static**: Yes, Wumpus and pits do not move.
- **Discrete**: Yes.
- **Single-agent**: Yes, Wumpus is essentially a part of the environment.
- **Fully observable**: No, only **local** perception.
  - This is our first example of partial observability.
- **Episodic**: No, what was observed before is very useful.

The agent needs to maintain a model of the world and to update this model upon percepts.

We will use **logical reasoning** to overcome the initial ignorance of the agent.

## Exploring the Wumpus world (1)

1,4	2,4	3,4	4,4
1,3	2,3	3,3	4,3
1,2	2,2	3,2	4,2
1,1 <b>A</b> OK	2,1 OK	3,1	4,1

**A** = Agent  
**B** = Breeze  
**G** = Glitter, Gold  
**OK** = Safe square  
**P** = Pit  
**S** = Stench  
**V** = Visited  
**W** = Wumpus

(a)

1,4	2,4	3,4	4,4
1,3	2,3	3,3	4,3
1,2	2,2 <b>P?</b>	3,2	4,2
1,1 <b>V</b> OK	2,1 <b>A</b> <b>B</b> OK	3,1 <b>P?</b>	4,1

(b)

(a) Percept = [None, None, None, None]

(b) Percept = [None, Breeze, None, None]

## Exploring the Wumpus world (2)

1,4	2,4	3,4	4,4
1,3 <b>W!</b>	2,3	3,3	4,3
1,2 <b>A</b> <b>S</b> <b>OK</b>	2,2  <b>OK</b>	3,2	4,2
1,1  <b>V</b> <b>OK</b>	2,1 <b>B</b>  <b>V</b> <b>OK</b>	3,1 <b>P!</b>	4,1

**A** = Agent  
**B** = Breeze  
**G** = Glitter, Gold  
**OK** = Safe square  
**P** = Pit  
**S** = Stench  
**V** = Visited  
**W** = Wumpus

1,4	2,4 <b>P?</b>	3,4	4,4
1,3 <b>W!</b>	2,3 <b>A</b> <b>S</b> <b>G</b> <b>B</b>	3,3 <b>P?</b>	4,3
1,2 <b>S</b> <b>V</b> <b>OK</b>	2,2  <b>V</b> <b>OK</b>	3,2	4,2
1,1  <b>V</b> <b>OK</b>	2,1 <b>B</b>  <b>V</b> <b>OK</b>	3,1 <b>P!</b>	4,1

(a)

(b)

(a) Percept = [Stench, None, None, None]

(b) Percept = [Stench, Breeze, Glitter, None]

# Logical agents

- Most useful in non-episodic, partially observable environments.
- Logic (knowledge-based) agents combine:
  - A **knowledge base** (**KB**): a list of facts that are known to the agent.
  - Current **percepts**.
- Hidden aspects of the current state are **inferred** using rules of inference.
- Logic provides a good formal language for both
  - Facts, encoded as **axioms**.
  - Rules of **inference**.

**function** KB-AGENT(*percept*) **returns an** *action*

**persistent:** *KB*, a knowledge base

*t*, a counter, initially 0, indicating time

TELL(*KB*, MAKE-PERCEPT-SENTENCE(*percept*, *t*))

*action*  $\leftarrow$  ASK(*KB*, MAKE-ACTION-QUERY(*t*))

TELL(*KB*, MAKE-ACTION-SENTENCE(*action*, *t*))

*t*  $\leftarrow t + 1$

**return** *action*

# Propositional logic

## Syntax

- The **syntax** of propositional logic defines allowable **sentences**.
- The syntax of propositional logic is formally defined by the following **grammar**:

$$\begin{aligned} \textit{Sentence} &\rightarrow \textit{AtomicSentence} \mid \textit{ComplexSentence} \\ \textit{AtomicSentence} &\rightarrow \textit{True} \mid \textit{False} \mid P \mid Q \mid R \mid \dots \\ \textit{ComplexSentence} &\rightarrow (\textit{Sentence}) \mid [\textit{Sentence}] \\ &\mid \neg \textit{Sentence} \\ &\mid \textit{Sentence} \wedge \textit{Sentence} \\ &\mid \textit{Sentence} \vee \textit{Sentence} \\ &\mid \textit{Sentence} \Rightarrow \textit{Sentence} \\ &\mid \textit{Sentence} \Leftrightarrow \textit{Sentence} \end{aligned}$$

OPERATOR PRECEDENCE :  $\neg, \wedge, \vee, \Rightarrow, \Leftrightarrow$

## Semantics

- In propositional logic, a **model** is an assignment of truth values for every proposition symbol.
  - E.g., if the sentences of the knowledge base make use of the symbols  $P_1, P_2$  and  $P_3$ , then one possible model is  $m = \{P_1 = \text{False}, P_2 = \text{True}, P_3 = \text{True}\}$ .
- The **semantics** for propositional logic specifies how to (recursively) evaluate the **truth value** of any complex sentence, with respect to a model  $m$ , as follows:
  - The truth value of a proposition symbol is specified in  $m$ .
  - $\neg P$  is true iff  $P$  is false;
  - $P \wedge Q$  is true iff  $P$  and  $Q$  are true;
  - $P \vee Q$  is true iff either  $P$  or  $Q$  is true;
  - $P \Rightarrow Q$  is true unless  $P$  is true and  $Q$  is false;
  - $P \Leftrightarrow Q$  is true iff  $P$  and  $Q$  are both true or both false.




# Wumpus world sentences

- Let  $P_{i,j}$  be true if there is a pit in  $[i, j]$ .
- Let  $B_{i,j}$  be true if there is a breeze in  $[i, j]$ .

Examples:

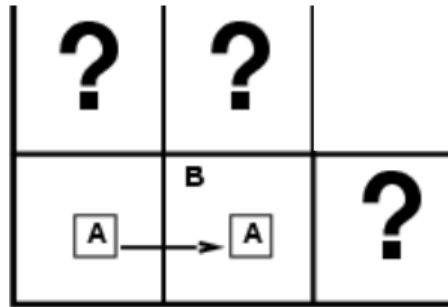
- There is no pit in  $[1, 1]$ :
  - $R_1 : \neg P_{1,1}$ .
- Pits cause breezes in adjacent squares:
  - $R_2 : B_{1,1} \Leftrightarrow (P_{1,2} \vee P_{2,1})$ .
  - $R_3 : B_{2,1} \Leftrightarrow (P_{1,1} \vee P_{2,2} \vee P_{3,1})$ .
  - These are true in all wumpus worlds.
- Breeze percept for the first two squares, for the specific world we consider:
  - $R_4 : \neg B_{1,1}$ .
  - $R_5 : B_{2,1}$ .



# Entailment

- We say a model  $m$  satisfies a sentence  $\alpha$  if  $\alpha$  is true in  $m$ .
- $M(\alpha)$  is the set of all models that satisfy  $\alpha$ .
- $\alpha \models \beta$  iff  $M(\alpha) \subseteq M(\beta)$ .
  - We say that the sentence  $\alpha$  entails the sentence  $\beta$ .
  - $\beta$  is true in all models where  $\alpha$  is true.
  - That is,  $\beta$  follows logically from  $\alpha$ .
- In other words, entailment enables logical inference.


# Wumpus models



- Let us consider possible models for KB assuming only pits and a reduced Wumpus world with only 5 squares and pits.
- We consider the situation after:
  - detecting nothing in  $[1, 1]$ ,
  - moving right, sensing breeze in  $[2, 1]$ .

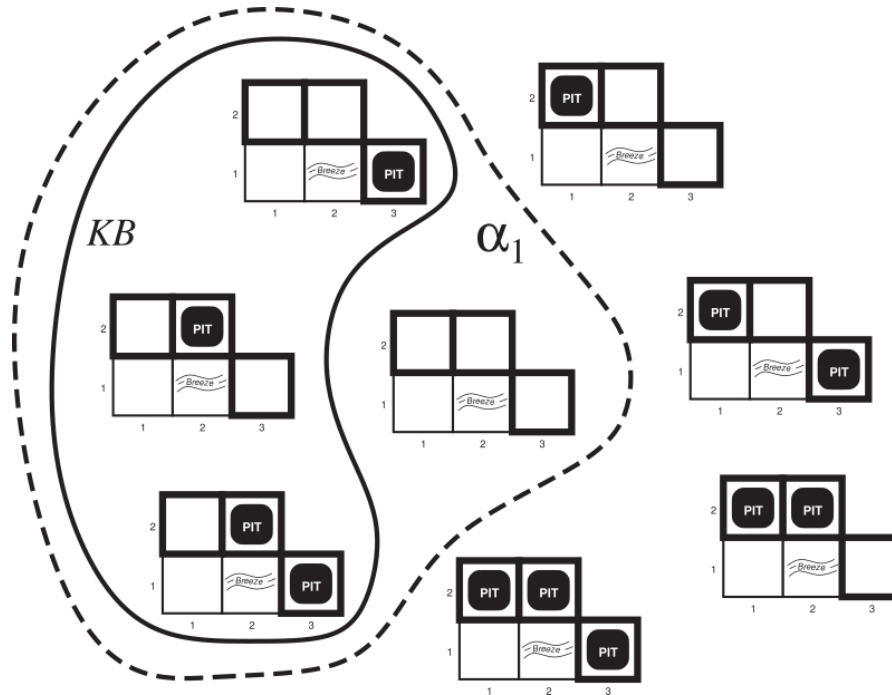
## Exercise

How many models are there?

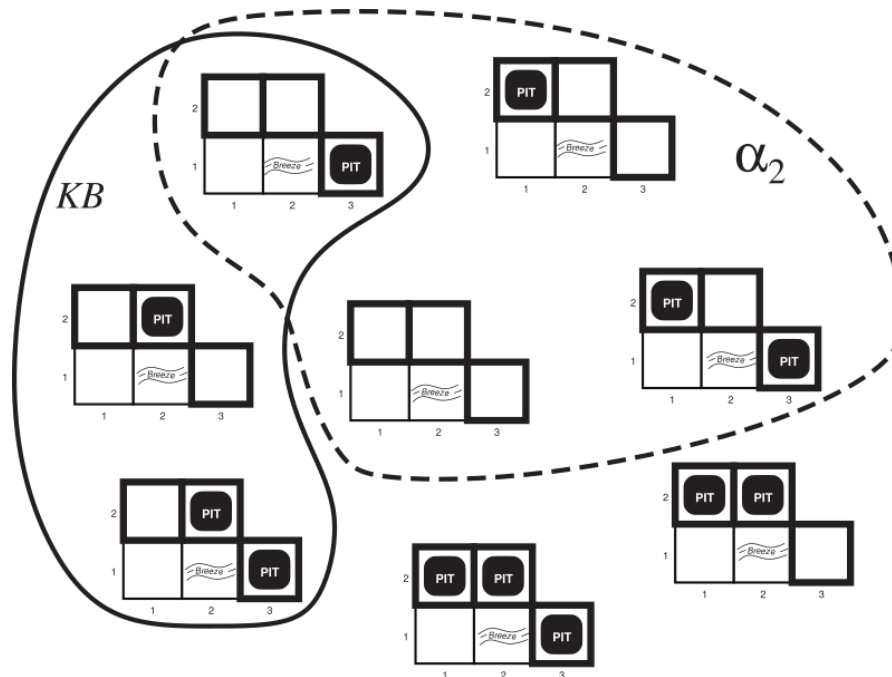


- All 8 possible models in the reduced Wumpus world.
- The knowledge base **KB** contains all possible Wumpus worlds consistent with the observations and the physics of the world.

# Entailments



- $\alpha_1 = "[1, 2]$  is safe". Does  $KB$  entails  $\alpha_1$ ?
- $KB \models \alpha_1$  since  $M(KB) \subseteq M(\alpha_1)$ .
  - This proof is called **model checking** because it **enumerates** all possible models to check whether  $\alpha_1$  is true in all models where  $KB$  is true.



- $\alpha_2 = "[2, 2]$  is safe". Does **KB** entails  $\alpha_2$ ?
- **KB**  $\not\models \alpha_2$  since  $M(\text{KB}) \not\subseteq M(\alpha_2)$ .
- We **cannot** conclude whether  $[2, 2]$  is safe (it may or may not).

# Unsatisfiability theorem

$\alpha \models \beta$  iff  $(\alpha \wedge \neg\beta)$  is unsatisfiable

- A sentence  $\gamma$  is unsatisfiable iff  $M(\gamma) = \{\}$ .
  - i.e., there is no assignment of truth values such that  $\gamma$  is true.
- Proving  $\alpha \models \beta$  by checking the unsatisfiability of  $\alpha \wedge \neg\beta$  corresponds to the proof technique of reductio ad absurdum.
- Checking the satisfiability of a sentence  $\gamma$  can be cast as CSP!
  - More efficient than enumerating all models, but remains NP-complete.
  - Alternatively, propositional satisfiability (SAT) solvers can be used instead of CSPs. These are tailored for this specific problem. Many of them are variants of backtracking.

# Limitations

- Representation of **informal** knowledge is difficult.
- Hard to define provable **plausible** reasoning.
- **Combinatorial explosion** (in time and space).
- Logical inference is only a part of intelligence.



# Summary

- Constraint satisfaction problems:
  - States are represented by a set of variable/value pairs.
  - Backtracking, a form of depth-first search, is commonly used for solving CSPs.
  - The complexity of solving a CSP is strongly related to the structure of its constraint graph.
- Logical agents:
  - Intelligent agents need knowledge about the world in order to reach good decisions.
  - Logical inference can be used as tool to reason about the world, in particular to infer parts that are not observable.
    - The inference problem can be cast as the problem of determining the unsatisfiability of a formula.
    - This in turn can be cast as a CSP.



# References

- Newell, A., & Simon, H. (1956). The logic theory machine--A complex information processing system. IRE Transactions on information theory, 2(3), 61-79.
- McCarthy, J. (1960). Programs with common sense (pp. 300-307). RLE and MIT computation center.